

AI for Financial Practitioners

Lecture 2 · Part II

Autonomy, Delegation and System Governance

How institutions delegate discretion without confusing capability with authority

The moment the assistant is allowed to act

The institutional problem changes when a system moves from producing information to altering the environment.

Assistance

- ▶ Retrieve evidence
- ▶ Calculate controlled metrics
- ▶ Draft analysis
- ▶ Propose an action



Delegated action

- ▶ Send an alert
- ▶ Update institutional state
- ▶ Open or route a case
- ▶ Move money or market exposure

Same model. Different authority. Different risk.

Capability is not authority

Technical capability asks *can it?*; institutional delegation asks *may it?*



Financial analogy

An analyst may be capable of trading without possessing trading authority. A lawyer may understand a contract without authority to bind the firm.

AI implication

API access, credentials or tool connectivity must never silently become delegated institutional authority.

Authority must be explicit, scoped and attributable

Autonomy = delegated discretion


Autonomy is a bounded grant over several dimensions—not an intrinsic property of a model.

Delegation dimension	Question for the institution
Objective	What outcome may the system pursue?
Action scope	Which actions are permitted or prohibited?
Information	Which data and sources may it use?
Resources	What money, exposure or compute may it consume?
Time	For how long and under what operating conditions?
Escalation	When must discretion contract back to a human?

The objective is not maximum autonomy. It is governable autonomy.

A ladder of bounded delegation

Authority can expand progressively—but more autonomy is not automatically more advanced.



Principle

Delegation should expand only when evidence supports the specific additional authority.

Consequence

A highly intelligent advisory system can rationally have less authority than a simple deterministic cash-sweep service.

Governance must be executable architecture

Policy becomes credible when controls sit in the action path—not in a document beside it.



Enforce at execution

Identity · permissions · amount limits · eligible instruments/accounts · policy version · duplicate exception path · owner · revocation capability protection

Preserve after execution

Confirmation · audit trail · reconciliation ·

Meaningful human oversight

A human in the loop is not a control if the human cannot realistically understand, challenge or stop the action.

Oversight requires

- ▶ Understandable evidence and rationale
- ▶ Sufficient decision time
- ▶ Genuine authority to reject or modify
- ▶ Manageable queue and case volume
- ▶ Clear escalation ownership

Warning signs

- ▶ Rubber-stamp approvals
- ▶ Reviewer overload during stress
- ▶ Automation bias / out-of-the-loop loss
- ▶ High confidence masking missing evidence
- ▶ Human approval after action is effectively irreversible

Oversight quality is an operating capacity, not a workflow label

Revocation, containment and controlled contraction

Autonomy is governable only if authority can contract when evidence, controls or operating assumptions deteriorate.



Triggers

Stale data, ambiguous governing document, limit breach, provider/model change, tool outage, reconciliation uncertainty.

Design test

Can the system continue to add value at a lower authority level rather than fail open or stop being useful?

Worked case: bounded treasury action

Treasury automation is credible when the action space is small, observable, enforced and reversible where possible.

Permitted

- ▶ Forecast daily cash positions
- ▶ Prepare transfers among approved internal accounts
- ▶ Execute below a defined threshold
- ▶ Operate during defined hours
- ▶ Reconcile confirmation to intended transfer

Not delegated

- ▶ Add beneficiaries
- ▶ Change thresholds
- ▶ Alter risk appetite or forecast objective
- ▶ Override stale critical data
- ▶ Blindly retry uncertain payments

Boundedness is not a limitation to apologize for. It is what makes delegation credible.

The board view: a portfolio of delegations

Senior governance should oversee autonomous services as a portfolio of authorities—not as isolated AI projects.

For every autonomous service, know:

- ▶ Business purpose and accountable owner
- ▶ Action classes and maximum consequence
- ▶ Evidence supporting the current delegation level
- ▶ Critical data, model, identity and tool dependencies
- ▶ Monitoring, escalation and revocation path

Executive principle

**Delegate only what you
can**

**define · observe · govern
· withdraw**

Next: adaptive institutions, recursive improvement and systems that participate in their own evolution.