

AI for Financial Practitioners

Lecture 1 · Part I

From Statistical Models to Generative Intelligence

A pedagogical monograph

How the language of finance expanded from estimation and prediction to representation, generation and the architecture of intelligent work

Alejandro Reynoso

Cambridge Judge Business School

Michaelmas 2026

ABSTRACT

This monograph develops a practical map of artificial intelligence for financial decision-makers. It distinguishes estimation, prediction, representation and generation; explains why a prediction is not a decision; and shows how context, retrieval, memory and deterministic tools convert a fluent model into a dependable working system. Its central argument is architectural: modern AI is most valuable when each component is assigned a bounded responsibility and when evidence, calculation, judgment and authority remain visible. A running credit-research example connects the conceptual distinctions to model selection, validation, human review and the design of institutionally credible workflows.

MASTER OF FINANCE
Michaelmas 2026

LECTURE 1 · PART I

From Statistical Models to Generative Intelligence

A visual guide to the key ideas of the monograph

UNIVERSITY OF
CAMBRIDGE
Judge Business School

1 FROM INFERENCE TO PREDICTION TO GENERATION
A century of models, each expanding what is possible in finance.

Statistical Models 1900s – 1970s	Machine Learning 1980s – 2000s	Representation Learning 2000s – 2010s	Generative Intelligence 2010s – today
 Inference Estimation Structure Hypothesis testing What is the relationship?	 Prediction Patterns Generalization Out-of-sample performance What is likely to happen?	 Understanding Embeddings Context Unstructured data How is the world represented?	 Generation Reasoning Composition Natural language interface What can be created or done?

2 PREDICTION IS NOT DECISION
A forecast estimates what may happen; a decision requires objectives, constraints, costs and authority.

A forecast describes the world.

A prediction is a probability distribution, not a recommendation.

A decision chooses an action.

- Objectives**
What are we trying to achieve?
- Constraints**
What are the limits?
- Costs**
What are the trade-offs?
- Uncertainty**
What might go wrong?
- Authority**
Who can decide and act?

Example: Portfolio Allocation

Forecast
Expected return: 8%
Volatility: 12%

Decision
Increase allocation by 5%

- ☒ Subject to risk limits
- ☒ Considering costs
- ☒ Within mandate
- ☒ Approved by the investment committee

3 GENERATIVE INTELLIGENCE CHANGES THE INTERFACE
Natural language becomes a powerful interface for analysis, synthesis, reasoning and creation.

Analysis Summarize, compare and extract insights from complex information.	Synthesis Combine information from multiple sources into coherent views.	Reasoning Work through complex questions with explicit assumptions and logic.	Creation Produce text, code, models, scenarios and structured outputs.
--	--	---	--

4 MODELS DO NOT OPERATE ALONE
A foundation model becomes truly useful when combined with context, tools, data and memory.

Context

- Task and objectives
- Domain and audience
- Constraints and requirements
- Institutional policies

Data and Retrieval

- Internal documents
- Market and reference data
- Real-time information
- Provenance and verification

Memory

- Conversation history
- Persistent knowledge
- User preferences
- Institutional memory

Tools

- Data analysis (Python, SQL)
- Modeling and computation
- External APIs and services
- Transaction and workflow systems

Deterministic Computation

- Calculations and simulations
- Optimization
- Risk and scenario analysis
- Validation and reconciliation

Governance and Guardrails

- Permissions and access control
- Policy and compliance checks
- Monitoring and auditability
- Human oversight and escalation

FOUNDATION MODEL
Language Reasoning Generation

5 FROM AI USER TO AI ARCHITECT
The critical professional transition is from asking for answers to designing the right combination of capabilities.

The AI User
Asks the model for answers

"Summarize this report."
"Analyze this data."
"What do you think?"

The AI Architect
Designs a system of capabilities

"Define the task and constraints."
"Design the architecture."
"Specify tools, data and governance."
"Deploy and monitor."

The professional opportunity in finance
The real advantage does not come from using a model once, but from designing and governing the right combination of models, data, tools, workflows and people for a specific financial task.

The purpose of this first half of Lecture 1 is not to replace the statistical language finance already knows. It is to place that language inside a larger map. The central question is simple: what changed when artificial intelligence moved from estimating relationships and predicting outcomes to representing complex information, generating new artifacts and participating in the organization of work?

Introduction: the familiar world and the new one

MAIN IDEA

Modern AI extends finance's long tradition of representing uncertainty; it expands the available operations without making statistical discipline obsolete.

For a financial practitioner, artificial intelligence can appear to have arrived all at once. A few years ago the normal conversation concerned regressions, factor models, credit scores, forecasts and optimization. Today the same practitioner can open a conversational interface, place an annual report in front of it, ask for an interpretation of the company, request a valuation framework, challenge the assumptions, generate code and obtain a draft investment memorandum before the first cup of coffee is finished. The experience is sufficiently different that it is tempting to describe everything that came before as “traditional” and everything that comes after as “AI.” That description is convenient, but intellectually unhelpful.

The better starting point is continuity. Finance has always built representations of an uncertain world. A balance sheet is a representation. A yield curve is a representation. A regression is a representation. A credit model is a representation. Each compresses reality so that a decision-maker can see something that would otherwise be difficult to see. Modern artificial intelligence enlarges the range of information that can be represented and the range of operations that can be performed on those representations. It does not repeal the older disciplines. It changes the architecture in which they can be combined.

Imagine an analyst studying a company whose margins have deteriorated. A classical statistical question might ask whether the deterioration is unusual relative to history or peers. A machine-learning question might ask which combination of variables best predicts future margin pressure. A representation-learning system might convert thousands of passages from filings, earnings calls and news into numerical representations

that preserve useful semantic relationships. A generative model might then synthesize those materials into competing explanations, draft questions for management, write code for a sensitivity analysis and explain the result in ordinary language. These are related activities, but they are not the same activity.

The practical objective of this monograph is therefore model literacy in a broader sense. The executive does not need to become a neural-network engineer. The executive does need to understand what kind of claim is being made when a system estimates, predicts, represents, generates, retrieves, calculates or recommends. Once those distinctions are clear, the apparent mystery of generative AI becomes much more manageable. We can ask a more useful question than “Is the model intelligent?” We can ask: “What responsibility are we giving this component, what evidence supports its output, and what other capabilities must surround it for the work to be dependable?”

1 From estimation to prediction to generation

MAIN IDEA

Estimation, prediction, representation and generation are complementary capabilities, each suited to a different type of financial question.

The natural point of departure for finance is statistical inference. Much of modern financial practice was built around the disciplined idea that observed data can be used to learn something about an underlying relationship. We estimate a beta, a default probability, an elasticity or the sensitivity of an option to a risk factor. The statistical model is valuable partly because it forces us to say what is being related to what. Even when the assumptions are imperfect, the structure is visible enough to be questioned.

Prediction changes the emphasis. Instead of asking primarily whether a coefficient has a particular interpretation, we may ask whether the model can forecast an outcome accurately on observations it has not seen. Machine learning greatly expands the set of functional relationships that can be fitted. Interactions need not be specified one by one. Nonlinearities can be learned. Large numbers of variables can be considered simultaneously. In credit, fraud detection, market microstructure and customer behavior, this can be extraordinarily useful.

But prediction is still a bounded operation. A model receives inputs and produces an

estimate of an outcome or a distribution over outcomes. It does not follow that the model understands the institutional purpose for which the prediction will be used. A probability of default is not a lending policy. A return forecast is not a portfolio. A volatility estimate is not a hedge. The prediction becomes economically meaningful only when it is connected to objectives, constraints and actions.

Representation learning adds another step. Much of the information that matters in finance does not arrive as a tidy matrix. It arrives as language, images, documents, transaction sequences and other high-dimensional objects. Modern neural systems can learn representations in which economically or semantically related objects become usefully close. This is one reason a system can search a large collection of documents for passages that are relevant in meaning even when they do not contain exactly the same words as the question.

Generative models then make a striking leap in the interface. Instead of returning only a score, class or forecast, they can produce language, code, images and structured outputs. They can transform one representation into another. A long filing can become a concise memo. A natural-language request can become a program. A table can become an explanation. A set of assumptions can become alternative scenarios. The model is no longer merely estimating a parameter hidden behind a specialist interface; it can participate in a dialogue through the same medium in which professionals describe problems to one another.

That is the source of much of the excitement. It is also the source of confusion. Generation feels general because language itself is general. Yet fluency should not be mistaken for universal reliability. A model that can discuss a discounted-cash-flow valuation may still make an arithmetic error in the calculation. A model that can summarize a covenant may retrieve the wrong version of the agreement. A model that can write plausible Python may produce code that executes while embodying the wrong economic definition. The expansion of capability is real, but it does not eliminate the need to assign responsibilities carefully.

A useful mental picture is not a ladder on which each new technology pushes the previous one away. It is a widening toolbox. Statistical inference remains useful when we care about explicit relationships and uncertainty. Predictive machine learning remains useful when out-of-sample performance is central. Representation learning becomes useful when the information is complex and unstructured. Generative models become useful when synthesis, transformation and interaction are required. A

mature financial architecture can use all of them.

2 Prediction is not decision

MAIN IDEA

A forecast describes a possible state of the world; a decision additionally requires objectives, constraints, trade-offs and legitimate authority.

The distinction between prediction and decision is one of the most important in the course because it protects us from a subtle category error. Suppose a model estimates that a bond has a 7 percent probability of default over a specified horizon. What should the portfolio manager do? The number alone cannot answer. The manager needs to know the expected recovery, the spread, the liquidity of the position, the correlations with the rest of the portfolio, the investment mandate, the relevant risk limits, the cost of trading and perhaps the consequences of being wrong in different directions.

In abstract form, a prediction describes something about the world. A decision chooses an action in that world. The bridge between them contains preferences and constraints. Two institutions can agree completely on the forecast and rationally choose different actions because they have different liabilities, capital positions, horizons or mandates. A pension fund and a hedge fund can observe the same expected return and volatility and still construct different portfolios. The disagreement need not be about the model. It may be about the objective.

This becomes especially important with generative AI because the conversational interface can collapse several stages into a single response. Ask, “Should we buy this stock?” and the system may produce a polished answer that blends facts, forecasts, valuation judgments and an apparent recommendation. The prose makes the transition feel natural. Institutionally, however, several distinct responsibilities have been compressed. What evidence was observed? Which quantities were calculated? Which future states were forecast? What objective function was assumed? What constraints were applied? Who has authority to act?

The discipline we want is not to prohibit integrated answers. It is to make their internal structure visible. Consider an investment memorandum. The first layer may contain observations: revenue grew by a certain amount; leverage increased; management

changed guidance. A second layer contains interpretations: perhaps the leverage increase reflects an acquisition rather than operating deterioration. A third layer may contain forecasts: cash conversion may normalize over the next four quarters. A fourth layer contains valuation and scenario analysis. Only then do we reach a proposed action, which must be evaluated against the portfolio mandate and the authority of the decision-maker.

The same distinction applies outside investing. A fraud model predicts the likelihood that a transaction is suspicious; the decision to block the transaction depends on customer impact, legal rules and escalation policy. A liquidity model forecasts cash needs; the treasury decision depends on funding costs, buffers and risk appetite. A regulatory model may estimate the probability that a submission contains a deficiency; the institution must still decide whether to file, revise or escalate.

This is why “better prediction” is not synonymous with “better institution.” A system can improve forecast accuracy and still worsen outcomes if it encourages excessive turnover, creates false confidence, ignores constraints or accelerates actions that should remain deliberate. Conversely, a modest model embedded in a disciplined decision process can create considerable value. The architecture around the prediction matters.

For the executive, this distinction changes the questions asked in a meeting. Instead of asking only, “How accurate is the model?” one asks, “What decision is this prediction intended to inform? What other information enters that decision? What are the costs of false positives and false negatives? What action is permitted at each confidence level?” Those questions move the conversation from model admiration to institutional design.

3 Generative intelligence changes the interface

MAIN IDEA

Natural-language interaction makes sophisticated analytical transformations accessible, but fluency must not be mistaken for factual or numerical reliability.

Generative AI matters in finance not merely because it can write. Its deeper significance is that natural language becomes a practical interface to a wide range of cognitive and computational operations. Professionals can express a task in the language of the business rather than first translating every task into the language of

software. That lowers the cost of experimentation and expands who can participate in building analytical workflows.

Consider the ordinary process of investigating a company. The analyst begins with a question that is not yet a model: “Why did free cash flow deteriorate despite reported earnings growth?” Traditionally, answering it may require searching several documents, locating footnotes, reconciling accounting definitions, constructing a spreadsheet and writing an explanation. A generative system can help coordinate these transformations. It can search semantically for relevant passages, propose a reconciliation, draft code to calculate the bridge, explain discrepancies and reformulate the analysis when the analyst challenges an assumption.

The remarkable feature is compositionality. The user does not need a separate interface for every micro-task. Language can specify the desired transformation. “Compare these two definitions.” “Use only information available before the earnings call.” “Show me the numerical bridge.” “Now write the strongest argument against our interpretation.” Each instruction modifies the work without requiring a new software application.

This creates a new role for professional judgment. The scarce skill is less often the ability to type a syntactically correct command and more often the ability to specify what a good answer means. If the analyst cannot distinguish reported EBITDA from an adjusted measure, the conversational interface will not rescue the analysis. If the credit officer does not know which covenant definition governs, a beautifully generated summary may simply make the error easier to circulate. Domain knowledge becomes more valuable because it defines the tests the system must pass.

Generative models also change the economics of prototypes. A team can test an idea before committing to a large software project. A compliance officer can describe a review process and see a first implementation. An investment team can experiment with a research assistant over a controlled document set. A CFO can ask for a scenario engine and inspect the code with technical colleagues. This does not mean the prototype is production-ready. It means the distance between idea and observable artifact has become much shorter.

That distinction is crucial. The first successful conversation with a model is evidence that a task may be tractable; it is not evidence that the task is operationally solved. The prototype has shown possibility. Production requires repeatability, data controls, permissions, evaluation, exception handling and accountability. In other words, the

new interface accelerates the beginning of the journey. It does not remove the rest of the journey.

There is also a cultural consequence. When language becomes an interface to computation, the boundary between “business” and “technology” becomes more permeable. A portfolio manager can participate more directly in defining analytical behavior. A technologist can inspect the business rationale expressed in ordinary language. The opportunity is not that everyone becomes a programmer. It is that the specification of intelligent work becomes a shared object that different specialists can examine together.

4 Models do not operate alone

MAIN IDEA

Dependable AI systems combine models with context, retrieval, memory, deterministic tools and human authorization in a deliberate division of labor.

A language model can appear self-contained because the conversation window hides much of the surrounding machinery. In a dependable financial application, however, the model is rarely the whole system. It sits inside an arrangement that supplies context, retrieves evidence, performs calculations, stores state and invokes tools. Understanding these surrounding capabilities is one of the simplest ways to become a more sophisticated user of AI.

Start with context. A model cannot respond appropriately to an institutional task unless the relevant circumstances are represented. “Is leverage high?” has no stable meaning without knowing the company, sector, accounting definition, period, currency and perhaps the mandate of the analyst. Context is the set of information that makes an instruction interpretable. Some context is supplied directly in the prompt; some comes from documents, databases, user identity, workflow state or prior interactions.

Retrieval addresses a different problem. The model may possess broad learned knowledge, but a financial institution often needs answers grounded in specific, current and authorized evidence. Retrieval brings relevant material into the working context: the latest filing, the approved policy, the executed contract, the research archive. Retrieval does not make an answer true by itself. It creates the possibility of

grounding the answer in evidence that can be inspected.

Memory introduces persistence. A single conversation may remember what was said earlier in that conversation. An institutional system may need something more deliberate: the status of a case, a prior approval, an unresolved exception, the assumptions used in yesterday's analysis. Memory is powerful precisely because it allows work to continue over time. It is also risky because stale, incorrect or unauthorized information can persist. The important question is not simply whether the system "has memory," but what is stored, for how long, with what provenance and under whose authority.

Tools allow the system to do things that language generation alone should not be trusted to do. A calculator can perform arithmetic. A pricing library can value an instrument under an explicit model. A database query can retrieve a position. A code interpreter can execute a sensitivity analysis. A workflow tool can prepare a ticket. The model may decide when and how to use the tool, but the tool provides a separate capability with its own controls and tests.

This division of labor is particularly useful in finance because many tasks combine interpretation with exactness. The model may be excellent at reading a paragraph and identifying that a debt covenant deserves attention. The exact covenant ratio should then be computed from controlled data using a transparent definition. The model may explain the result, compare it with the agreement and draft the escalation. We should not force one component to perform every responsibility merely because it can produce an answer that looks complete.

The architecture therefore becomes a composition. A generative model can orchestrate language and reasoning; retrieval can supply evidence; deterministic code can enforce calculations; specialized predictive models can estimate probabilities; memory can preserve approved state; humans can authorize consequential actions. Reliability emerges from the arrangement, not from a mystical property of a single model.

This is also why model comparisons, while useful, can become disproportionately important in public discussion. The institution does not deploy a benchmark score. It deploys a process. A somewhat stronger model inside a weak process can be less dependable than a somewhat weaker model inside a well-designed one. The executive question is therefore architectural: which capability should carry which responsibility?

5 From AI user to AI architect

MAIN IDEA

The mature practitioner moves from asking what a model can do to deciding which capability should carry each responsibility in a governed workflow.

The final idea of Part I is a change in posture. The novice asks, “What can this model do for me?” The more mature practitioner asks, “How should intelligence be assembled for this task?” The first question is about consumption. The second is about architecture.

Suppose an asset manager wants to monitor five hundred companies for deteriorating credit conditions. It would be easy to begin by asking a language model to read filings and produce alerts. The architect begins differently. What counts as deterioration? Which observations must be current? Which calculations must be exact? Which sources are authoritative? Which judgments can be probabilistic? Which outputs may remain drafts? Which events require human escalation? What evidence should accompany an alert? How will the institution know whether the service is improving or merely producing more prose?

These questions decompose the assignment into responsibilities. Filing ingestion is not the same as entity resolution. Entity resolution is not the same as ratio calculation. Ratio calculation is not the same as interpretation. Interpretation is not the same as forecasting. Forecasting is not the same as deciding whether to reduce exposure. Once the responsibilities are visible, technologies can be selected more intelligently.

This posture also makes experimentation safer. Instead of demanding a fully autonomous system on day one, the institution can begin with a bounded task. Perhaps the system only prepares a draft memorandum with citations and numerical checks. Analysts compare it with their own work. Failure cases are recorded. The system may then be allowed to monitor a larger universe, but consequential recommendations still require review. Authority can expand only when evidence justifies expansion.

The architect therefore thinks in terms of interfaces and tests. What information crosses from one component to another? What assumptions are preserved? What happens when evidence is missing? What is the baseline against which the new method is evaluated? A successful demonstration is not defined by surprise or eloquence. It is defined by whether the system performs a specified responsibility under conditions that matter to the institution.

This way of thinking is liberating because it removes the false choice between enthusiasm and skepticism. One does not need to believe that generative AI is either magical or useless. It is a powerful new class of capability. Its value depends on the work we ask it to perform and the architecture in which we place it. The task of leadership is to understand enough about the capabilities to design that architecture responsibly.

6 A taxonomy of claims for financial work

MAIN IDEA

Descriptive, inferential, predictive, generative and prescriptive claims require different evidence and must not inherit authority from one another.

The distinctions developed so far become easier to use when they are translated into a taxonomy of claims. Every output in an analytical workflow should be understood as making one or more of five claims: a descriptive claim about what has been observed, an inferential claim about a relationship, a predictive claim about an uncertain future, a generative claim about a useful transformation, or a prescriptive claim about an action. The categories can coexist in one memorandum, but they should not inherit one another's authority.

A descriptive claim might state that net debt increased by £400 million. Its quality depends on entity identity, accounting definition, date, currency and source. An inferential claim might argue that the increase is associated with an acquisition rather than a deterioration in operating cash flow. Its quality depends on a causal or interpretive argument. A predictive claim might estimate the probability that leverage will breach a threshold. Its quality depends on calibration and performance under relevant conditions. A generative claim might compress the evidence into a readable paragraph. Its quality depends on faithfulness and usefulness. A prescriptive claim might recommend reducing exposure. Its quality depends on objectives, constraints, alternatives and delegated authority.

This taxonomy prevents a common transfer of confidence. A passage can be an excellent summary and a poor forecast. A forecast can be statistically strong and economically irrelevant. An economically attractive action can be prohibited by mandate. The fact that all three appear in the same fluent answer does not make them equally supported. The practitioner should therefore ask, sentence by sentence, what

kind of claim is being made and what evidence would justify it.

The older statistical literature is especially useful here because it treats assumptions and generalization as objects of scrutiny rather than invisible properties of an answer. Bishop's account of probabilistic modeling and the treatment of model assessment by Hastie, Tibshirani and Friedman remain relevant even when the interface is generative [1, 2]. The new interface changes how work is requested and communicated; it does not abolish the need to distinguish a fitted pattern from a dependable claim.

An institution can make the taxonomy operational. Drafts may label observations with source and date, calculations with definitions, forecasts with horizons and uncertainty, interpretations with competing explanations, and recommendations with the constraints assumed. The result need not read like a legal document. A compact evidence table or a set of margin labels is enough to preserve the internal structure. What matters is that the reader can see where retrieval ends, where calculation begins, where judgment enters and where authority would be required.

7 Representation, embeddings and retrieval in practice

MAIN IDEA

Semantic representations improve access to unstructured evidence, but financial retrieval still requires metadata, version control and inspectable provenance.

Representation learning deserves a closer examination because it supports many capabilities that appear magical at the interface. Traditional information retrieval often begins with matching words. That works well when the vocabulary of the query matches the vocabulary of the document. It works less well when a user asks about “funding pressure” and the relevant passages discuss deposit attrition, wholesale refinancing, collateral calls or a reduction in committed facilities. A learned representation can place these passages near one another because their use in language suggests a semantic relationship.

An embedding is a numerical representation of an object such as a passage, table description or query. Distance in the representation space can serve as a rough measure of relevance. The approach is powerful because the same retrieval layer can search heterogeneous documents without requiring every synonym to be anticipated. It is limited because semantic proximity is not the same as legal, accounting or economic

equivalence. “Adjusted EBITDA” can be close in meaning across documents while differing materially in definition. A system can retrieve a thematically relevant passage and still retrieve the wrong definition for the calculation at hand.

The design implication is that retrieval should combine several signals. Semantic similarity can produce candidates. Metadata can restrict the candidates to the correct entity, period, document type and authorization class. Keyword or exact-match rules can preserve identifiers, defined terms and clause references. A reranking step can compare the question with the candidate passages in more detail. The final answer should preserve citations that allow a person or a downstream control to inspect the evidence.

Transformers made it possible to learn and operate on rich contextual representations at scale [4]; modern deep-learning systems generalize this representational approach across language, images and sequences [3]. Yet a financial retrieval service is not validated by demonstrating that it can find an impressive passage. It should be tested on a corpus that contains near-duplicates, amendments, superseded policies, multiple currencies, similar issuer names and questions for which no authorized answer exists.

Two error rates matter. Retrieval recall asks whether the relevant evidence was included among the candidates. Retrieval precision asks how much irrelevant or misleading material was included. Increasing the number of retrieved passages may improve recall while filling the working context with distractors. The best operating point depends on the downstream responsibility. A brainstorming assistant can tolerate broad retrieval. A covenant-monitoring service may require strict document lineage and explicit abstention when the governing agreement is unavailable.

The practitioner should also distinguish retrieval from memory. Retrieval usually searches an external collection for evidence relevant to the current task. Memory preserves selected state from prior work. If an analyst approved a definition of recurring revenue for one company, storing that definition may prevent repeated work. But it may also become stale after an acquisition or accounting change. Stored state therefore needs an owner, a timestamp, provenance and a rule for revalidation. Intelligence becomes dependable when representations make information accessible without erasing the institutional conditions under which the information is valid.

8 Evaluation beyond a persuasive demonstration

MAIN IDEA

Deployment evidence must test a clearly defined responsibility across ordinary cases, consequential edge cases, abstentions, costs and operational failures.

A persuasive demonstration normally answers the easiest question: can the system produce a useful result at least once? Deployment requires a harder question: how often does it perform the specified responsibility under the conditions that matter, and what happens when it does not? Foundation models acquire broad capabilities from scale, but their broad competence and broad risk surface must be considered together [5].

Evaluation should begin with a task definition. “Analyze companies” is not testable. “Identify material changes in liquidity risk from newly published issuer documents and produce a cited draft for analyst review within thirty minutes” is closer. It specifies an event, a responsibility, an output, an audience and a service expectation. Acceptance criteria can then be divided into evidence quality, numerical correctness, analytical usefulness, operational performance and control behavior.

Evidence quality asks whether every material factual claim is supported by an authorized source and whether the source is the correct version. Numerical correctness asks whether quantities reconcile to controlled data and approved definitions. Analytical usefulness asks whether the draft identifies the issues that change the decision rather than merely restating the document. Operational performance includes latency, cost, availability and the burden placed on reviewers. Control behavior asks whether the system abstains, escalates and records its work when evidence is missing or contradictory.

A useful evaluation set is not a random pile of old examples. It should represent the distribution of work and deliberately over-represent consequential edge cases. For credit monitoring, these might include restated filings, pro-forma acquisitions, covenant amendments, revolving-credit drawdowns, entities with similar names, foreign-language supplements and periods of market stress. Each example should have an expected evidence set and a documented reason why a particular conclusion or escalation is appropriate.

The baseline also matters. If an analyst currently spends two hours producing a

strong memorandum, the new system should be compared with that process, not with perfection in the abstract. It may create value by reducing preparation time while leaving the final decision with the analyst. Conversely, a faster draft that doubles verification time has not necessarily improved the process. The unit of evaluation should be useful completion: the complete cost and quality of getting the responsibility to an acceptable state.

Evaluation continues after release. The institution should record corrections, abstentions, source failures, reviewer overrides and downstream outcomes. These data reveal where the design envelope is too broad or the controls too weak. They do not automatically authorize the system to learn from every correction. Some corrections reflect preference; others reveal a factual error; still others reflect a policy change. The institution must classify them before deciding whether to alter prompts, retrieval, code, training data or workflow authority.

9 Worked case: a credit-monitoring architecture

MAIN IDEA

A credible credit-monitoring service assigns identity, retrieval, calculation, prediction, synthesis and approval to components designed for those tasks.

Consider a service that monitors a portfolio of corporate bonds. Its purpose is not to decide trades. Its bounded responsibility is to detect new evidence of credit deterioration, calculate a controlled set of metrics and prepare an analyst-ready briefing. The architecture begins with an event: a new filing, earnings-call transcript, rating action or material company announcement enters an approved source channel.

The first component resolves identity. It confirms the issuer, parent, guarantors, security and relevant reporting period. This step may seem administrative, but an error here contaminates everything that follows. The second component classifies the document and checks whether it supersedes an earlier version. The third performs hybrid retrieval: semantic search finds passages on liquidity, refinancing, covenants and operating pressure, while exact rules locate financial tables, defined terms and dates.

A deterministic calculation service then constructs metrics such as leverage, interest coverage, free cash flow and near-term maturities. Each metric carries its formula,

units, source cells and reconciliation status. The language model does not invent the number. It interprets the controlled result in the context of the retrieved evidence. A separate predictive model may estimate transition or default risk, but its horizon, training population and calibration date remain attached to the estimate.

The generative component assembles the briefing. A useful structure is: what changed; why it may matter; what evidence supports the interpretation; what competing explanation remains plausible; which calculations were performed; what information is missing; and what question the analyst should answer next. The draft contains citations to the source passages and links each reported metric to its calculation record. If the evidence is inconsistent, the system generates an exception rather than resolving the inconsistency rhetorically.

The analyst remains responsible for the investment judgment. The service can prioritize the case, prepare alternative scenarios and calculate sensitivities. It cannot infer the portfolio's complete objective from the issuer evidence. The analyst brings exposure, liquidity, mandate, relative value and knowledge of other positions. If the system eventually receives authority to create a watch-list entry automatically, that is a separate delegation decision supported by evidence about false positives, missed events and the reversibility of the action.

This case illustrates why “using a language model” is an incomplete description. The service contains source controls, entity resolution, retrieval, deterministic calculation, perhaps a specialized predictive model, generation, human review and workflow state. Removing any one component changes the claim the service can responsibly make. The architecture is not overhead surrounding the intelligence. It is the means by which different forms of intelligence become institutionally useful.

10 A practical design exercise for executives

MAIN IDEA

Executives can design AI responsibly by mapping claims, evidence, capabilities, action boundaries and acceptance tests before choosing a technology.

A leadership team can apply the chapter without selecting a vendor or writing code. Choose one recurring responsibility that currently consumes professional attention. Write the desired output in one sentence, identify its user and state the consequence

of error. Then list the observations required, the interpretations made, the calculations performed, the forecasts introduced and the actions that follow.

For each item, assign a capability. Authoritative data systems should preserve facts. Deterministic tools should carry exact calculations. Predictive models should provide estimates whose validation matches the population and horizon. Generative models should synthesize, translate and interact. Humans should retain objectives, contested judgments and authority where consequences justify it. The allocation will not be perfect, but the exercise exposes responsibilities that were previously hidden inside a job title or a polished answer.

Next, define the evidence boundary. Which sources are authorized? How current must they be? What metadata must travel with them? What should happen when the evidence is absent? Define the action boundary separately. Which outputs are drafts, which may update internal state, and which could create an external consequence? A system that is safe for drafting may not be safe for sending, and a system that is safe for sending information may not be safe for placing an order.

Finally, specify a test. Select ordinary cases, difficult cases and cases that should force abstention. Compare the complete workflow with the present baseline. Record not only answer quality but review time, correction type, latency, cost and the frequency with which a human must reconstruct the evidence. The exercise changes the conversation from whether AI is impressive to whether a particular allocation of capability improves a defined responsibility.

11 Seminar application: dissecting an AI equity note

MAIN IDEA

Reconstructing a polished equity note into observations, calculations, forecasts and recommendations reveals the architecture hidden by fluent prose.

Imagine that an equity-research assistant produces a two-page note after a company reports results. The note states that revenue exceeded expectations, margins disappointed, management guidance became more cautious and the shares appear undervalued. It includes a valuation table and recommends buying the stock. The note is articulate, timely and directionally plausible. The seminar task is to reconstruct the architecture hidden inside the prose.

Begin with observation. Identify every statement that should be traceable to a source: reported revenue, consensus expectations, margin, guidance, share price, balance-sheet items and management quotations. For each observation, specify entity, period, unit, source and observation time. Mark any value whose definition may differ across the filing, data vendor and analyst model.

Next isolate calculation. Rebuild growth, margin, net debt, enterprise value and valuation multiples using controlled formulas. Ask whether the assistant used reported or adjusted measures and whether the peer set shares those definitions. If the table cannot be reproduced, the problem is not merely presentation; the economic conclusion lacks a dependable numerical foundation.

Then identify interpretation and prediction. “Management became more cautious” is an interpretation that should point to changed language or numbers. “Margins will recover” is a forecast requiring a horizon, mechanism and competing scenario. “Undervalued” combines forecast, discount rate, terminal assumptions and comparison with price. The group should state which assumptions drive most of the conclusion and what new evidence would falsify them.

Finally examine prescription. A buy recommendation requires portfolio context that may not appear in the company note: mandate, current exposure, liquidity, risk budget, alternatives and timing. The same evidence can support different actions for a concentrated long-only fund, a market-neutral strategy and an institution prohibited from owning the security. The model’s recommendation has silently assumed an objective.

Redesign the service so that it produces an analyst-ready draft rather than an apparently self-sufficient answer. Assign source retrieval, numerical calculation, scenario construction, prose generation and approval to appropriate components. Define three cases in which it must abstain or escalate. Select an evaluation baseline and measure both note quality and verification time. The exercise is complete when the group can explain not whether the model is intelligent, but which claims the complete system can responsibly make.

Closing synthesis: from capability to process

MAIN IDEA

The central shift is from evaluating an impressive model in isolation to designing a complete process in which evidence, intelligence and authority remain visible.

We can now see the first half of Lecture 1 as a movement across five ideas. Finance begins with a familiar discipline of estimation and prediction. Modern AI expands the space through representation learning and generation. The conversational interface makes complex transformations accessible, but it can also conceal the difference between evidence, inference and action. Context, retrieval, memory and tools reveal that useful intelligence is composed rather than monolithic. The professional therefore moves from being a user of an impressive model to becoming an architect of a task.

The natural next question is temporal. Once the capabilities have been allocated, how do they operate together continuously? A one-off memorandum can be inspected before anyone acts. A service that watches hundreds of companies cannot rely on the same informal safeguards. New information arrives. Old information becomes stale. Actions fail. Outcomes appear later. Permissions differ across users and tasks. The architecture must therefore become a process.

Part II begins at exactly that boundary. We will follow intelligence through a governed cycle of perception, judgment, planning, execution and feedback. The emphasis will shift from what individual components can do to how an institution can make their combined behavior dependable. That shift—from capability to operating system—is the bridge from generative intelligence to institutional intelligence.

References

MAIN IDEA

These sources provide the statistical, machine-learning and foundation-model concepts that support the monograph's architectural framework.

References

- [1] Christopher M. Bishop. *Pattern Recognition and Machine Learning*. Springer, 2006.
- [2] Trevor Hastie, Robert Tibshirani and Jerome Friedman. *The Elements of Statistical Learning*, 2nd ed. Springer, 2009.
- [3] Ian Goodfellow, Yoshua Bengio and Aaron Courville. *Deep Learning*. MIT Press, 2016.
- [4] Ashish Vaswani et al. "Attention Is All You Need." *Advances in Neural Information Processing Systems*, 30, 2017.
- [5] Rishi Bommasani et al. "On the Opportunities and Risks of Foundation Models." Stanford Center for Research on Foundation Models, 2021.